

SOHAN SHINGADE

sohanshingade@gmail.com | 650 305 6043 | github.com/sohan-shingade | U.S. Citizen

EDUCATION

University of California San Diego — B.S. in Data Science Sep 2023 – Present (Expected June 2027)
Coursework: Machine Learning, Deep Learning, Statistical Learning, Econometrics, Time Series Analysis, Distributed Systems, Algorithms & Data Structures, Databases & SQL, Optimization, Computer Vision, NLP, Cloud Computing / HPC

EXPERIENCE

Triton Trading Group Sep 2025 – Present
Quantitative Research Lead — La Jolla, CA

- Lead a **12+ member** quant research team operating a profitable live portfolio with a realized **Sharpe ratio of 3.1**; manage signal ideation through walk-forward backtesting to live execution
- Built AI-assisted research pipelines for feature selection, regime filtering, and automated parameter sweeps; increased team iteration throughput by **2.5x**
- Developed transaction-cost-aware backtesting infrastructure with slippage modeling and MEV-aware execution analysis; top strategy achieved **2.0+ Sharpe** with <8% max drawdown OOS
- Standardized evaluation protocols (turnover constraints, factor exposure checks, leakage tests) and review gates; reduced overfit submissions by **35%**, improved paper-trade promotion rate by **40%**
- Produced weekly research reports synthesizing market trends, strategy performance, and actionable recommendations; reduced manual status-gathering by **60%** through automated dashboards

Yonder Dynamics Sep 2024 – Present
Software Engineer — La Jolla, CA

- Developed and deployed PyTorch ML models for real-time inference on an autonomous robotics platform; achieved **92% accuracy at 30 FPS** on edge hardware via batch optimization
- Built distributed Python/C++ services for data ingestion, model serving, and telemetry via REST APIs; improved pipeline throughput by **1.8x** via async I/O and batching
- Built a React/TypeScript dashboard for real-time system monitoring; reduced operator debugging time by **45%** with interactive filtering and log replay
- Containerized services with Docker and orchestrated CI/CD pipelines; cut environment setup time from **2 hours to 15 minutes** for new contributors

Shepherd Ventures Sep 2024 – Present
Data Platform Engineer — San Diego, CA

- Built Python+SQL data pipelines on AWS (S3, EC2, Lambda) to ingest and serve datasets from 5+ sources; reduced data refresh latency by **46%** and eliminated manual ETL
- Trained and evaluated ML models (scikit-learn, XGBoost, PyTorch) with rigorous validation; improved experiment cycle time by **2.0x** via modular feature engineering
- Designed a PostgreSQL analytics platform with parameterized endpoints and Redis caching; improved load times by **3x** for internal consumers
- Developed NLP sentiment parsing on financial news and social data; reduced false positives by **30%** through calibration and cross-signal research

Stockbee Jun 2024 – Aug 2024
Software Engineer Intern — Cupertino, CA

- Designed and shipped production backend services in Go and Python for real-time trading analytics; reduced **p95 API latency by 35%** via profiling and connection pooling
- Built automated data validation and regression testing gates; blocked **15k+ malformed records** and cut production incidents by **22%**
- Implemented structured logging, Grafana dashboards, and on-call runbooks across 3 microservices; reduced MTTR by **40%** with end-to-end tracing
- Designed scalable REST API endpoints with pagination, filtering, and rate limiting; served **50k+ daily requests** at <100ms median latency

PROJECTS

Flint — Solana DeFi Algorithmic Trading & Backtesting Engine [Open Source](#)
Local-first algo trading framework for Solana DeFi. 13 data providers, event-driven backtest engine with realistic fill/fee models, Optuna hyperparameter optimization, paper trading with risk guards. React 19 UI with TradingView charts. Python/FastAPI/DuckDB. 497 tests.

Jobslop — VC-Backed Startup Job Aggregator [Live](#)
Full-stack app aggregating 12,800+ listings from 35+ VC portfolio boards (Sequoia, a16z, Greylock). Async Python scrapers, Turso DB, Next.js/SSR. AI resume matching via Groq/Llama 3.3 70B. Sub-200ms loads, 99.8% uptime.

Order Book Simulator — Price-Time Priority Matching Engine [C++](#)
Full limit order book with per-price FIFO queues, deterministic matching, partial fills, event-level audit logs. Microstructure features (spread, depth, imbalance), 500k+ events/sec.

Wildfire Risk Heatmap — Geospatial ML Pipeline
Binary classification on 2.4M rows from 6 data sources. scikit-learn with extreme class imbalance handling, ROC/PR AUC evaluation, interactive folium heatmaps.

TECHNICAL SKILLS

Languages: Python, Rust, TypeScript/JavaScript, SQL, C++, Go, Java, Kotlin, VBA, Bash, R
ML/AI: PyTorch, scikit-learn, XGBoost, LLM integration (Groq, LangChain), model serving, inference optimization, NLP
Backend & Infra: FastAPI, Node.js, REST/WebSocket APIs, PostgreSQL, Redis, DuckDB, Docker, AWS (S3/EC2/Lambda), CI/CD, Kubernetes
Data & Finance: pandas, NumPy, ETL pipelines, batch processing, financial modeling, backtesting, portfolio analytics, market microstructure